

# Reinforcement Learning

## Lecture 2

Georg Martius

Distributed Intelligence / Autonomous Learning Group, Uni Tübingen, Germany

October 22, 2024

EBERHARD KARLS  
UNIVERSITÄT  
TÜBINGEN



## Tutor sessions:

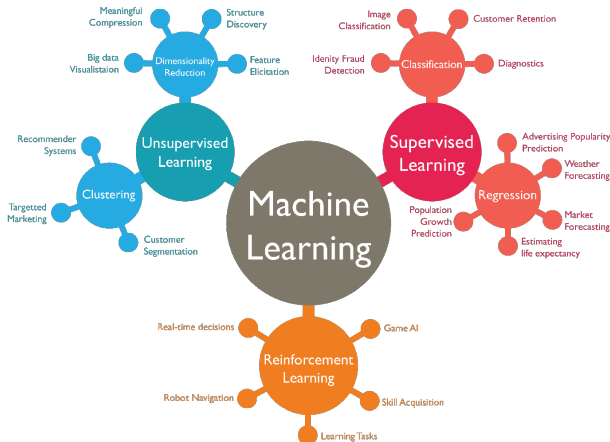
- ▶ Hörsaal N11 (next door) (change in room in 3 weeks)
- ▶ Hörsaal N03 (Hörsaalzentrum Morgenstelle) (change in room in 3 weeks)
- ▶ Hörsaal TTR2 (Obere Viehweide, Maria-von-Linden-Str. 6, Tübingen AI center)

## Lectures:

- ▶ Today's lecture will go over-time (would like to cover until Value-Iteration for homeworks)
- ▶ Next week: I am on a scientific Conference (EWRL). Dr. René Geist will substitute me


# Machine Learning Overview (Reminder)

Three different classes of tasks:



RL is special, because: interaction with a system, generation of own data, sequential

# Agent and Environment (Reminder)



# The Reinforcement Learning Problem (Reminder)

Behave such that maximal **expected (discounted) future return** is achieved


- ▶ Behave: find **Policy** that determines actions
- ▶ Optimal w.r.t. expected return: **Value function**
- ▶ Maybe **Model** the environment

Two fundamental problems in sequential decision making

- ▶ Reinforcement Learning:
  - ▶ The environment is initially unknown
  - ▶ The agent interacts with the environment
  - ▶ The agent improves its policy
  - ▶ a.k.a. learning by doing, trial and error learning
- ▶ Planning:
  - ▶ A model of the environment is known
  - ▶ The agent performs computations with its model (without any external interaction)
  - ▶ The agent improves its policy
  - ▶ a.k.a. deliberation, reasoning, introspection, pondering, thought, search

[Slide adapted from David Silver]

# Atari Example: Reinforcement Learning




- ▶ Rules of the game are **unknown**
- ▶ Learn directly from interactive game-play
- ▶ Pick actions on joystick, see pixels and scores

[Slide adapted from David Silver]

# Atari Example: Planning

- ▶ Rules of the game are **known**
- ▶ Can query emulator  
perfect model inside agent's brain
- ▶ If I take action  $a$  from state  $s$ :
  - ▶ what would the next state be?
  - ▶ what would the score be?
- ▶ Plan ahead to find optimal policy  
e.g. tree search



[Slide adapted from David Silver]



# Relationship between Planning and Reinforcement Learning

- ▶ Planning: **on-the-fly** computation of **best action**
  - ▶ typically short-horizon optimization only
- ▶ RL: learns (for a long time) to find the **best policy**
  - ▶ solves the global optimization problem
  - ▶ **amortizes previous interactions** into a policy ➡ fast at runtime
- ▶ Is planning and RL mutually exclusive?
  - ▶ No, but traditionally treated by different communities
  - ▶ can be combined (model-based RL, AlphaGo, etc)

Subproblems within Reinforcement Learning:

- ▶ **Prediction**: evaluate the future  
Given a policy: How good is the agent?
- ▶ **Control**: optimize the future  
Find the best policy with respect to current knowledge

# Markov Decision Processes

Formally describe environments for reinforcement learning

A Markov process is a memoryless random process, i.e. a sequence of random states  $S_1, S_2, \dots$  with the Markov property.

## Reminder: Markov property

A state  $S_t$  is Markov if and only if

$$P(S_{t+1} \mid S_t) = P(S_{t+1} \mid S_1, \dots, S_t)$$

## Definition (Markov Process/ Markov Chain)

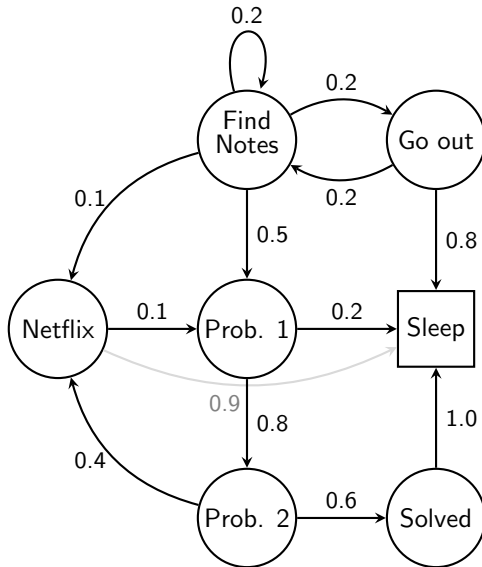
A *Markov Process* (or *Markov Chain*) is a tuple  $(\mathcal{S}, \mathcal{P})$

- ▶  $\mathcal{S}$  is a (finite) set of states
- ▶  $\mathcal{P}$  is a state transition probability matrix,


$$\mathcal{P}_{ss'} = P(S_{t+1} = s' \mid S_t = s)$$

## Example: Student Markov Chain

Let's consider the evening of a student in this class ;-)



## Example: Student Markov Chain Episodes



Sample episodes for starting from  $s_1 = \text{FN}$

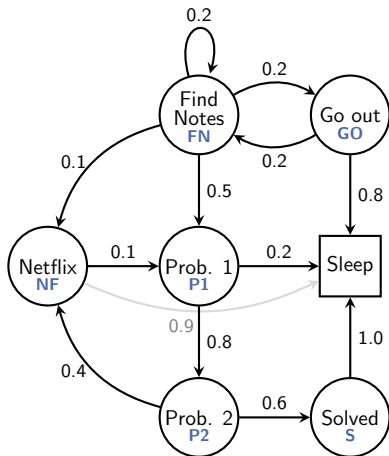
$s_1, s_2, \dots, s_T$

- ▶ FN P1 P2 S Sleep
- ▶ FN GO FN FN P1 Sleep
- ▶ FN P1 P2 NF P1 P2 S Sleep
- ▶ FN GO FN P1 P2 NF P1 P2 NF P1 Sleep

# Example: Student Markov Chain Episodes

Transition Matrix:

$$\mathcal{P}_{ss'} = P(S_{t+1} = s' \mid S_t = s)$$



Transition Matrix:

$$\mathcal{P} = \begin{matrix} & \begin{matrix} \text{FN} & \text{P1} & \text{P2} & \text{S} & \text{GO} & \text{NF} & \text{Sleep} \end{matrix} \\ \begin{matrix} \text{FN} \\ \text{P1} \\ \text{P2} \\ \text{S} \\ \text{GO} \\ \text{NF} \\ \text{Sleep} \end{matrix} & \begin{bmatrix} 0.2 & 0.5 & & & 0.2 & 0.1 & \\ & & 0.8 & & & & 0.2 \\ & & & 0.6 & & 0.4 & \\ & & & & & & 1.0 \\ 0.2 & & & & & & 0.8 \\ & 0.1 & & & & & 0.9 \\ & & & & & & 1 \end{bmatrix} \end{matrix}$$

A Markov reward process is a Markov chain with values.

## Definition (MRP)

A *Markov Reward Process* is a tuple  $(\mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma)$

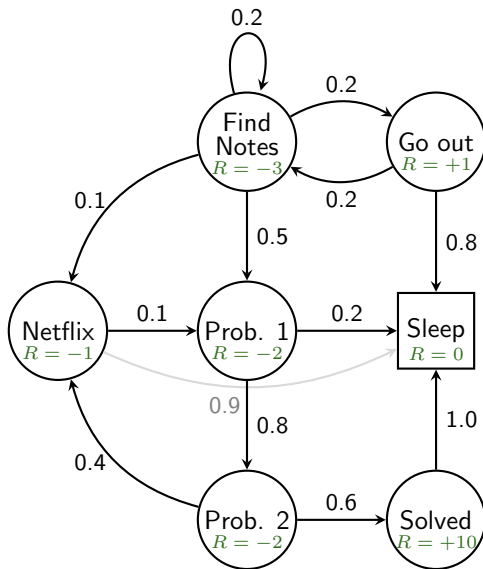
- ▶  $\mathcal{S}$  is a finite set of states
- ▶  $\mathcal{P}$  is a state transition probability matrix,

$$P(\mathcal{S}_{t+1} \mid \mathcal{S}_t) = P(\mathcal{S}_{t+1} \mid \mathcal{S}_1, \dots, \mathcal{S}_t)$$

- ▶  $\mathcal{R}$  is a reward function,  $\mathcal{R}_s = \mathbb{E}[R_{t+1} \mid \mathcal{S}_t = s]$
- ▶  $\gamma$  is a discount factor,  $\gamma \in [0, 1]$



## Example: Student MRP



## Definition

The *return*  $G_t$  is the total discounted reward from time-step  $t$ .

$$G_t = R_{t+1} + \gamma R_{t+2} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- ▶ The discount  $\gamma \in [0, 1]$  is the present value of future rewards
- ▶ The value of receiving reward  $R$  after  $k + 1$  time-steps is  $\gamma^k R$ .
- ▶ This values immediate reward above delayed reward.
  - ▶  $\gamma$  close to 0 leads to “myopic” evaluation
  - ▶  $\gamma$  close to 1 leads to “far-sighted” evaluation

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Why is discounting often used?

- ▶ Mathematically convenient to discount rewards (keeps returns finite)
- ▶ A way to model the uncertainty about the future (since the model may not be exact)
- ▶ Animal/human behaviour shows preference for immediate rewards

In some cases *undiscounted* Markov reward processes (i.e.  $\gamma = 1$ ), are considered, e.g. if all sequences terminate.

The value function describes the value of a state (in the stationary state)


## Definition

The state *value function*  $v(s)$  of an MRP is the expected return starting from state  $s$

$$v(s) = \mathbb{E}[G_t \mid S_t = s]$$


## Example: Value Function for Student MRP

Extremely myopic  $\gamma = 0$ :  $v(s)$




## Example: Value Function for Student MRP

Now more farsighted  $\gamma = 0.9$ :  $v(s)$



## Example: Value Function for Student MRP

Now fully farsighted  $\gamma = 1.0$ :  $v(s)$



# Bellman Equation (MRP) I

Idea: Make value computation recursive by tearing apart contributions from:

- ▶ immediate reward
- ▶ and from discounted future rewards

$$\begin{aligned}v(s) &= \mathbb{E}[G_t \mid S_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) \mid S_t = s]\end{aligned}$$

Mh... need Expectation over  $S_{t+1}$


Use transition matrix to get probabilities of succeeding state:

$$v(s) = \mathcal{R}_s + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'} v(s')$$



## Example: Bellman Equation for Student MRP

$\gamma = 1.0$ :  $v(P2)$



$$v(P2) = 3.64 = -2 + 0.4 \cdot (-0.91) + 0.6 \cdot 10$$

# Bellman Equation (MRP) II

Bellman equations in matrix form:

$$v = \mathcal{R} + \gamma \mathcal{P}v$$

where  $v \in \mathbb{R}^{|S|}$  and  $\mathcal{R}$  are vectors

The Bellman equation can be solved explicitly (in closed form):

$$v = (\mathbb{I} - \gamma \mathcal{P})^{-1} \mathcal{R}$$

► computational complexity is  $O(|S|^3)$

# Markov Decision Process

A Markov reward process has no agent, there is no influence on the system. An MRP with an active agent forms a Markov Decision Process.

- ▶ Agent takes **decision** by executing **actions**
- ▶ State is Markovian

## Definition (MDP)

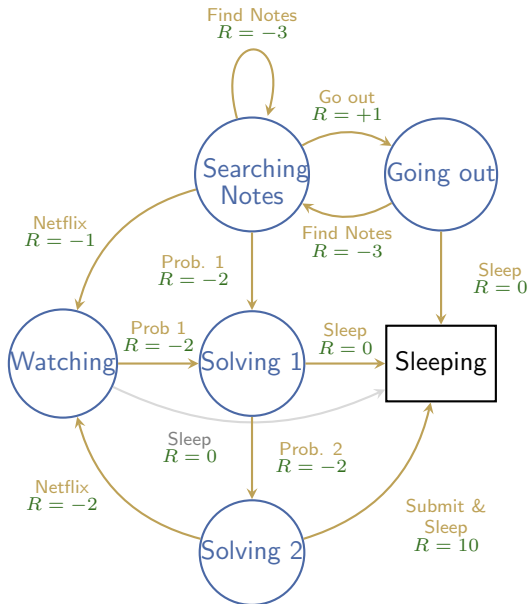
A **Markov Decision Process** is a tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$

- ▶  $\mathcal{S}$  is a finite set of states
- ▶  $\mathcal{A}$  is a finite set of actions
- ▶  $\mathcal{P}$  is a state transition probability matrix,

$$\mathcal{P}_{ss'}^a = P(S_{t+1} = s' \mid S_t = s, A_t = a)$$

- ▶  $\mathcal{R}$  is a reward function,  $\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$
- ▶  $\gamma$  is a discount factor,  $\gamma \in [0, 1]$

# Example: Student MDP



# How to model decision taking?

The agent has an action function called **policy**.

## Definition

A **policy**  $\pi$  is a distribution over actions given states,

$$\pi(a \mid s) = P(A_t = a \mid S_t = s)$$

- ▶ Since it is a Markov process the policy only depends on the current state
- ▶ Implication: policies are stationary (independent of time)

An MDP with a given policy turns into a MRP:

$$\mathcal{P}_{ss'}^\pi = \sum_{a \in \mathcal{A}} \pi(a \mid s) \mathcal{P}_{ss'}^a,$$

$$\mathcal{R}_s^\pi = \sum_{a \in \mathcal{A}} \pi(a \mid s) \mathcal{R}_s^a$$

# Modelling expected returns in MDP

How good is each state when we follow the policy  $\pi$ ?

## Definition

The **state-value function**  $v_\pi(s)$  of an MDP is the expected return when starting from state  $s$  and following policy  $\pi$ .

$$v_\pi(s) = \mathbb{E}[G_t \mid S_t = s]$$

Should we change the policy?

How much does choosing a different action change the value?

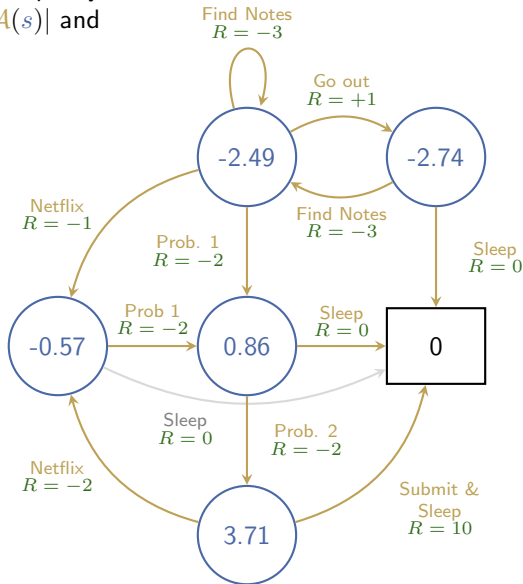
## Definition

The **action-value function**  $q_\pi(s, a)$  of an MDP is the expected return when starting from state  $s$ , taking action  $a$ , and then following policy  $\pi$ .

$$q_\pi(s, a) = \mathbb{E}[G_t \mid S_t = s, A_t = a]$$

# Example: State-Value function for Student MDP

$v_{\pi}(s)$  for uniform policy  
 $\pi(a | s) = 1/|\mathcal{A}(s)|$  and  
 $\gamma = 1$



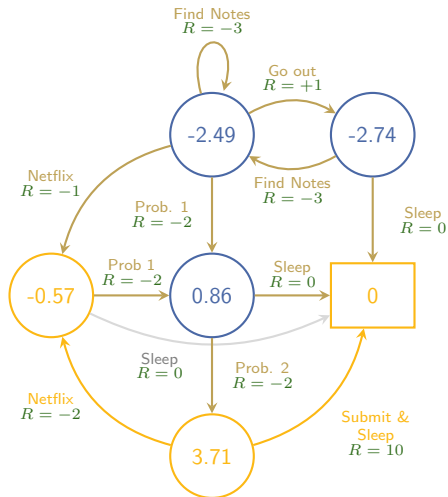
# Bellman Expectation Equation

Recall:

**Bellman Equation:** decompose expected reward into **immediate reward** plus **discounted value of successor state**:

$$v_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s]$$

$$\begin{aligned} v_{\pi}(P2) &= 3.71 \\ &= 0.5 \cdot (-2 - 0.57) + 0.5 \cdot (10 + 0) \end{aligned}$$



$$\pi(a \mid s) = 1/|\mathcal{A}(s)| \text{ and } \gamma = 1$$



# Bellman Expectation Equation

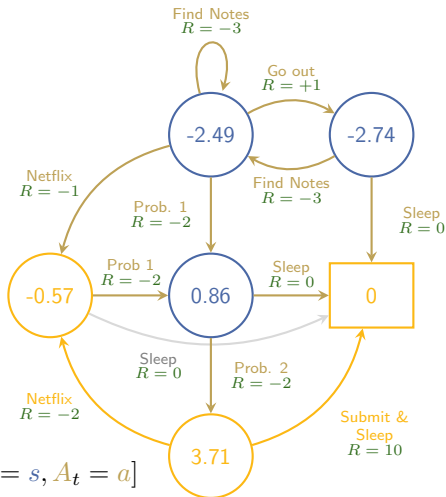
Recall:

**Bellman Equation:** decompose expected reward into **immediate reward** plus **discounted value** of **successor state**:

$$v_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s]$$

The **action-value function** can be similarly decomposed,

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a]$$



$$\pi(a \mid s) = 1/|\mathcal{A}(s)| \text{ and } \gamma = 1$$

# Bellman Equation: update of $v_\pi$

Value function can be derived from  $q_\pi$ :


$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a | s) q_\pi(s, a)$$

... and  $q$  can be computed from transition model

$$q_\pi(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_\pi(s')$$

Substituting  $q$  in  $v$ :

$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a | s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_\pi(s') \right)$$



Since a policy induces an MRP  $v_\pi$  can be directly computed (as before)

$$v = (\mathbb{I} - \gamma \mathcal{P}^\pi)^{-1} \mathcal{R}^\pi$$

But do we want  $v_\pi$ ?

We want to find the **optimal** policy and its value function!

## Definition

The **optimal state-value function**  $v_*(s)$  is the maximum value function over all policies

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

## Definition

The **optimal action-value function**  $q_*(s, a)$  is the maximum action-value function over all policies

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

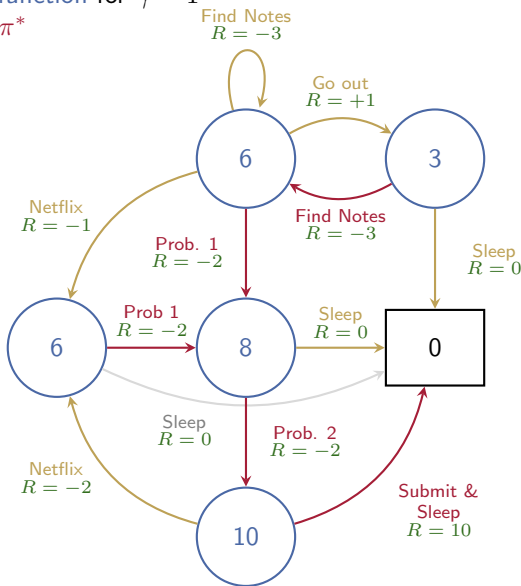
What does it mean?

- ▶  $v_*$  specifies the best possible performance in an MDP
- ▶ Knowing  $v_*$  solves the MDP (how? we will see...)

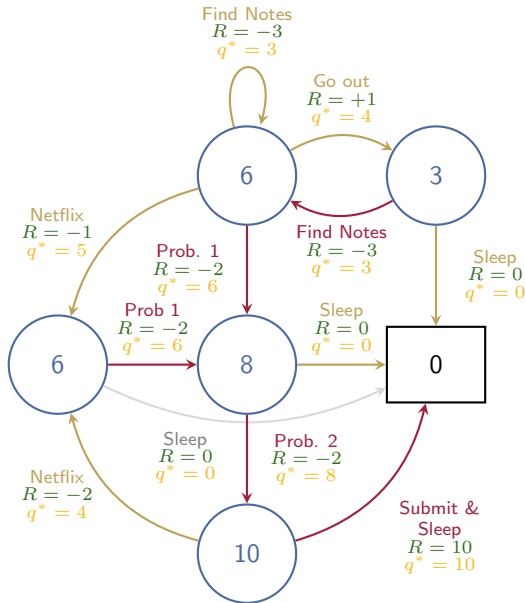
# Example: Optimal Value Function $v_*$ in Student MDP

Optimal value function for  $\gamma = 1$

optimal policy  $\pi^*$



# Example: Optimal State Function $q_*$ in Student MDP



To **solve** an MDP, we need to obtain the **optimal policy**.  
Define a partial ordering over policies

$$\pi \geq \pi' \text{ if } v_\pi(s) \geq v_{\pi'}(s), \forall s$$

## Theorem

*For any Markov Decision Process*

- ▶ *There exists an optimal policy  $\pi_*$  that is better than or equal to all other policies,  $\pi_* \geq \pi, \forall \pi$*
- ▶ *All optimal policies achieve the optimal state-value function,  $v_{\pi_*}(s) = v_*(s)$*
- ▶ *All optimal policies achieve the optimal action-value function,  $q_{\pi_*}(s, a) = q_*(s, a)$*

# Finding an Optimal Policy

Given the optimal action-value function  $q_*$ :  
How do we get the optimal policy?

Discuss with your peers ...

The optimal policy is given by **maximizing**  $q_*$

$$\pi_*(a \mid s) = \mathbb{I}[a = \arg \max_{a \in \mathcal{A}} q_*(s, a)]$$

$\mathbb{I}[\cdot]$  is Iverson bracket: 1 if *true*, otherwise 0.

- ▶ There is always a deterministic optimal policy for any MDP
- ▶ If we know  $q_*(s, a)$ , we immediately have the optimal policy (greedy)



# Bellman Equation for optimal value functions

Also for the optimal value functions we can use Bellmans optimality equations:  
Exercise!

Remember:


$$v_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a \mid s) q_{\pi}(s, a)$$

$$\pi_*(a \mid s) = \llbracket a = \arg \max_{a \in \mathcal{A}} q_*(s, a) \rrbracket$$

Bellman Optimality Equation is non-linear (because of max operation)

- ▶ No closed form solution (in general)
- ▶ Many iterative solution methods
  - ▶ Value Iteration
  - ▶ Policy Iteration
  - ▶ Q-learning
  - ▶ SARSA

# Questions?



[Image: [globalrobots.com](https://globalrobots.com)]

# Break



# Dynamic Programming

# Dynamic Programming?

**Dynamic:** sequential or temporal component of the problem


**Programming:** optimizing a “program”, i.e. a policy  
name like in linear programming  
(mathematical programming = optimization)

- ▶ A method for solving complex problems
- ▶ By breaking them down into subproblems
  - ▶ Solve the subproblems
  - ▶ Combine solutions to subproblems

# When can we use Dynamic Programming?

When problems have two properties:

- ▶ Optimal substructure
  - ▶ **Principle of optimality** applies
  - ▶ Optimal solution can be decomposed into subproblems
- ▶ Overlapping subproblems
  - ▶ Subproblems recur many times
  - ▶ Solutions can be cached and reused



# When can we use Dynamic Programming?

When problems have two properties:

- ▶ Optimal substructure
  - ▶ Principle of optimality applies
  - ▶ Optimal solution can be decomposed into subproblems
- ▶ Overlapping subproblems
  - ▶ Subproblems recur many times
  - ▶ Solutions can be cached and reused

Markov decision processes? ➡ Satisfy both properties!


- ▶ Bellman equation gives recursive decomposition
- ▶ Value function stores and reuses solutions



# Dynamic Programming to solve MDPs

Dynamic programming assumes full knowledge of the MDP!

Example MDP:



- ▶ Undiscounted episodic MDP ( $\gamma = 1$ )
- ▶ One terminal state (shown twice as shaded squares)
- ▶ Actions leading out of the grid leave state unchanged
- ▶ Reward is  $-1$  until the terminal state is reached

Can be used for **planning** in an MDP.

**Input:** MDP  $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$  and policy  $\pi$

## For prediction

Yields value function  $v_\pi$

	↕↕↕	↕↕↕	↕↕↕	
↕↕↕	↕↕↕	↕↕↕	↕↕↕	
↕↕↕	↕↕↕	↕↕↕	↕↕↕	
↕↕↕	↕↕↕	↕↕↕	↕↕↕	

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0


## For control

Yields optimal value function  $v_*$  and  $\pi_*$

	←	←	↖
↑	↖	↖	↓
↑	↖	↖	↓
↖	→	→	

# Prediction: Policy Evaluation

**Problem:** evaluate policy (find its value function)



**Solution:** Dynamic Programming to compute value function

Iterative algorithm to obtain:  $v_1 \rightarrow v_2 \rightarrow v_3 \rightarrow v_\pi$

## Iterative Policy Evaluation

- ▶ At each iteration  $k + 1$
- ▶ For all states  $s \in \mathcal{S}$
- ▶ Update  $v_{k+1}(s)$  from  $v_k(s')$  for all  $s'$
- ▶ where  $s'$  is a successor state of  $s$



# Iterative Policy Evaluation

Bellmann Expectation Equation:

$$v_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_k(s') \right)$$

$$v_{k+1}(s) = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi v^k$$

“Backup” diagram:



# Iterative Policy Evaluation: Small Gridworld

$v_k$  for the  
Random Policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

# What to do with the value function?

## Improve the Policy!

### Policy Iteration

Given a policy  $\pi$

1. Evaluate the policy  $\pi$  (Policy Evaluation)


$$v_{\pi}(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \dots \mid S_t = s]$$

2. Improve the policy by acting greedily with respect to  $v_{\pi}$  (Policy Improvement)

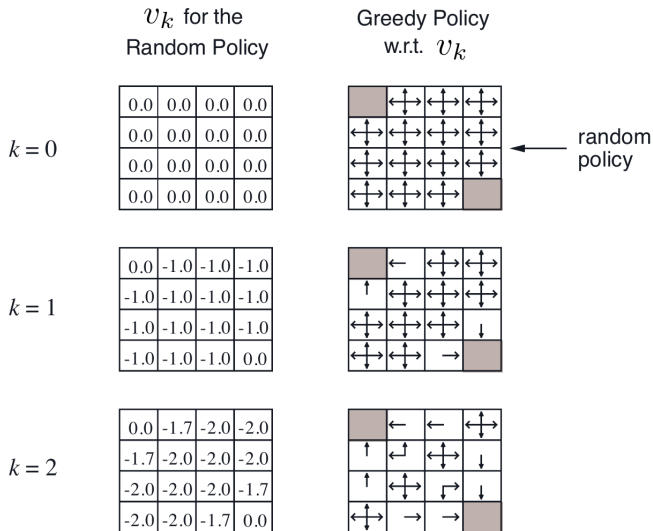
$$\pi' = \text{greedy}(v_{\pi})$$

3. iterate until policy does not change


► Policy iteration always converges to  $\pi^*$



# Policy Iteration in Small Gridworld (one iteration)



# Policy Iteration in Small Gridworld (one iteration)



- In Small Gridworld: Policy improvement converges after one iteration ( $\pi' = \pi^*$ )

- ▶ Consider a deterministic policy,  $a = \pi(s)$
- ▶ Policy *improvement* step by acting greedily

$$\pi'(s) = \arg \max_{a \in \mathcal{A}} q_{\pi}(s, a)$$

- ▶ This improves the value from any state  $s$  over one step,

$$q_{\pi}(s, \pi'(s)) = \max_{a \in \mathcal{A}} q_{\pi}(s, a) \geq q_{\pi}(s, \pi(s)) = v_{\pi}(s)$$

- ▶ It therefore improves the value function:  $v_{\pi'}(s) \geq v_{\pi}(s)$

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)) = \mathbb{E}_{\pi'} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s] \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) \mid S_t = s] \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 q_{\pi}(S_{t+2}, \pi'(S_{t+2})) \mid S_t = s] \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \dots \mid S_t = s] = v_{\pi'}(s) \end{aligned}$$



- ▶ If improvements stop,

$$q_{\pi}(s, \pi'(s)) = \max_{a \in \mathcal{A}} q_{\pi}(s, a) = q_{\pi}(s, \pi(s)) = v_{\pi}(s)$$

- ▶ Then the Bellman optimality equation has been satisfied

$$v_{\pi}(s) = \max_{a \in \mathcal{A}} q_{\pi}(s, a)$$

- ▶ Therefore  $v_{\pi}(s) = v_{*}(s)$  for all  $s \in \mathcal{S}$

➡  $\pi$  is an optimal policy




# Value Iteration

**Goal:** compute optimal value function and optimal policy

**Algorithm:** Like **Policy Iteration**, but with 1-step policy evaluation

1. **Policy Evaluation:** estimate  $v_\pi$  (not to convergence)
  2. **Policy Improvement:** generate  $\pi'$  with  $v_{\pi'} \leq v_\pi$
  3. iterate until value function does not change
- ▶ **Intuition:** start with final rewards and work backwards
  - ▶ given an optimal solution for some states, e.g.  $v_*(s)$ 
    - ➡ solution  $v_*(s)$  can be found by one-step lookahead

“Backup” diagram:



# Value Iteration

Improve subproblem solution  $v_k(s')$

based on current best estimates of  $v_*(s)$

## Value Iteration

1. For all states  $s \in \mathcal{S}$
2. Update  $v_{k+1}(s)$  from  $v_k(s')$

$$v_{k+1}(s) \leftarrow \max_{a \in \mathcal{A}} \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_k(s') \right) \quad \text{Bellman Optimality Eq.}$$

3. if value function changes goto (1): next iteration with  $k \leftarrow k + 1$
- ▶ There is not explicit policy here
  - ▶ Intermediate value functions might not correspond to any policy
  - ▶ Converges to  $v_*$
  - ▶ Is the same as **Policy Iteration**, but with **1-step Policy evaluation**

# Value Iteration: Example Shortest Path

g			

Problem

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$V_1$

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1

$V_2$

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-2
-2	-2	-2	-2

$V_3$

0	-1	-2	-3
-1	-2	-3	-3
-2	-3	-3	-3
-3	-3	-3	-3

$V_4$

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-4
-3	-4	-4	-4

$V_5$

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-5

$V_6$

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

$V_7$

# Synchronous Dynamic Programming Algorithms

Bellmann Expectation and Optimality Eq:

$$v_{k+1}(s) = \mathbb{E}_{a \sim \pi} \left( \mathcal{R}_s^a + \gamma \sum_{s' \in S} \mathcal{P}_{ss'}^a v_k(s') \right)$$
$$v_{k+1}(s) = \max_a \left( \mathcal{R}_s^a + \gamma \sum_{s' \in S} \mathcal{P}_{ss'}^a v_k(s') \right)$$

So far we looked at synchronous Methods:  
all states are updated at once.

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

Complexity: ( $m$  actions,  $n$  states)

- ▶ Using state-value function  $v(s)$ :  $O(mn^2)$  per iteration
- ▶ Using state-action-value function  $q(s, a)$ :  $O(m^2n^2)$  per iteration

- ▶ Model-free Prediction and Control