

# Reinforcement Learning Homework 4

Deadline 12th November 2024

Students Sahiti Chebolu, Surabhi S Nath, Xin Sui

## 1 Q-Learning and SARSA

(a)

Q-learning is considered an off-policy control method because it enables learning about policy  $\pi$  from experience sampled with  $\mu$  (as opposed to just from  $\pi$  as in SARSA). It uses the best local next action  $A'$  for value backup instead of the  $Q(S', A')$  in SARSA:

$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma \max_{a'} Q(S', a') - Q(S, A)) \quad (1)$$

(b)

If Q-learning action selection is greedy, this doesn't make it equivalent to SARSA. In SARSA, action selection (and the policy being improved) is  $\epsilon$ -greedy.

There's still the on vs. off policy distinction in terms of value backup/learning. They generally do not make the same action selections and weight updates because they learn from different Q values.

(c)

(a) Action going right should always be taken at state A because it results in on average higher rewards (0 as opposed to -0.1).

(b) We'd expect Q-learning to yield an optimal policy of going left at state A instead. This is because Q-learning uses the best next Q for value backup and is therefore 'biased' towards the best possible final reward ( $> 0$ ) instead of the average final reward (0).

## 2 Hands-on in Gridworld and Q-Learning

(a)

See Fig 1 below. Q learning predicted higher expected returns in general because it uses the best local next action  $A'$  for value backup, therefore it is (at least largely) unaffected by noise. However, there's an exception with the state in the upper left corner - Q learning produced smaller expected return there. This is probably a result of the  $\epsilon$ -greedy policy in Q learning as compared to the greedy policy in value iteration.

To make the values closer to optimal values, we need to decay  $\epsilon$  towards 0 and train for more episode. The optimal values should also be produced in an environment with 0 noise instead of the default 0.2.

(b)

See Fig 2 below. Here, the learned q-values are very small compared to results from value iteration. This is a result of the  $\epsilon$ -greedy policy in Q learning in addition to the highly negative reward associated with falling off the bridge. Even for 10000 episodes (Fig 3), the values obtained for value-iteration and Q-learning are wildly different. This is because Q-learning agent's value estimates are swayed by the large, negative rewards obtained randomly due to the  $\epsilon$ -greedy policy.

(c)

The value estimate for the start state from 300 Q-learning iterations is 2.74 (Fig 4). The average returns from these episodes is -26.88. For comparison, the value estimate from value iteration is -4.14 (Fig 5). The average returns is -2.95. The discrepancy in the Q-learning case is because the action selection is  $\epsilon$ -greedy, leading to negative reward of -100 sometimes due to random actions (instead of greedy action selection)

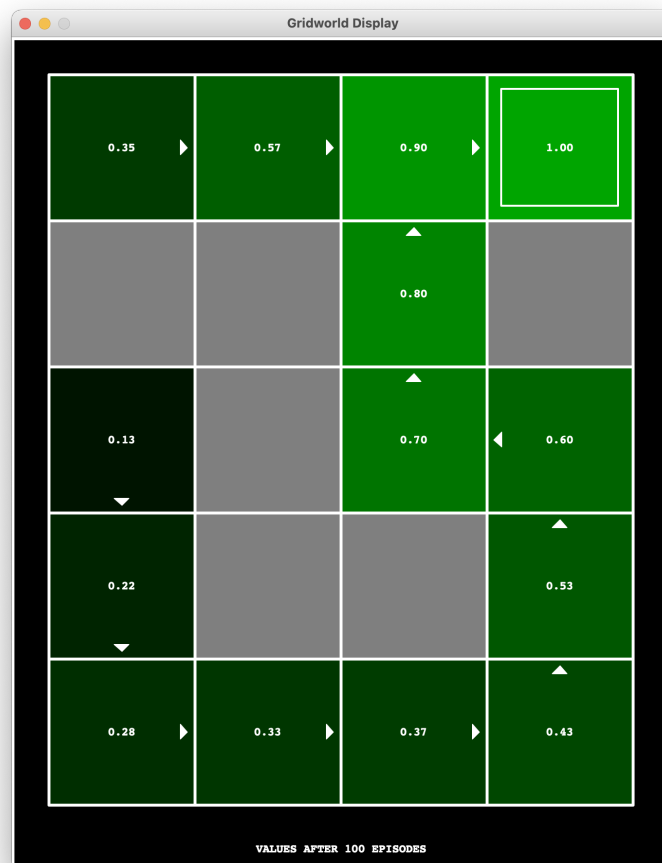


Figure 1: Value function obtained with Q learning with default parameters on MazeGrid after 100 episodes

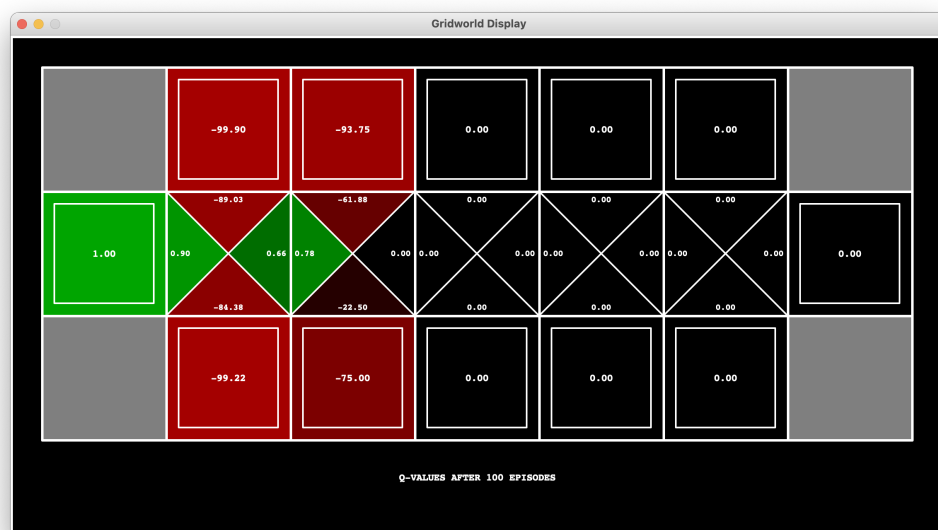


Figure 2: Q-function obtained with Q learning on BridgeGrid without noise after 100 episodes

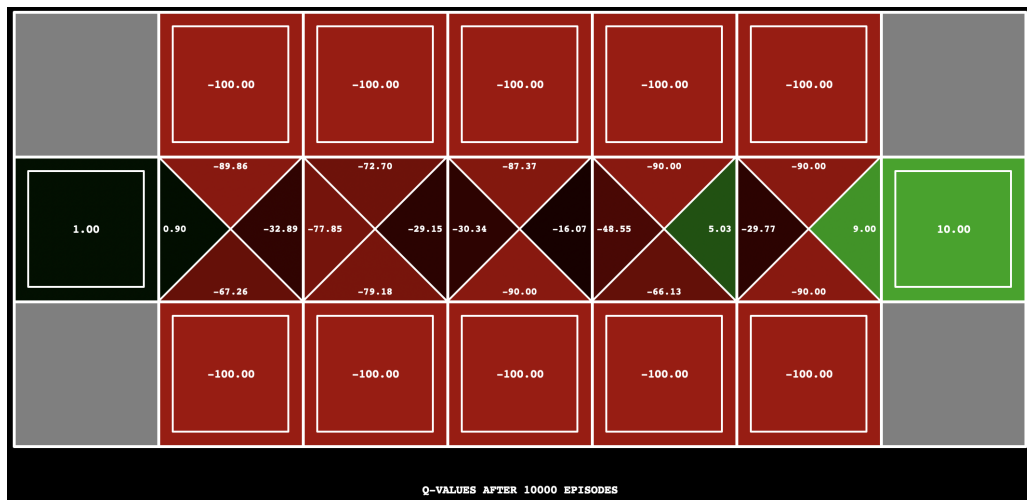


Figure 3: Q-function obtained with Q learning on BridgeGrid without noise after 10000 episodes




Figure 4: Value estimates obtained with Q learning on CliffGrid after 300 episodes




Figure 5: Value estimates obtained with value iteration on CliffGrid after 300 iterations

(d)

We show value estimates for book grid with 300 episodes of Q-learning (Fig 6) and 300 iterations of value iteration (Fig 7). The Q-learning agent is sensitive to the terminal state with negative reward (-1) obtained due to sub-optimal actions of the  $\epsilon$ -greedy policy, hence leading to avoidant behavior around the negative state. This persists even after 10000 episodes (Fig 8).

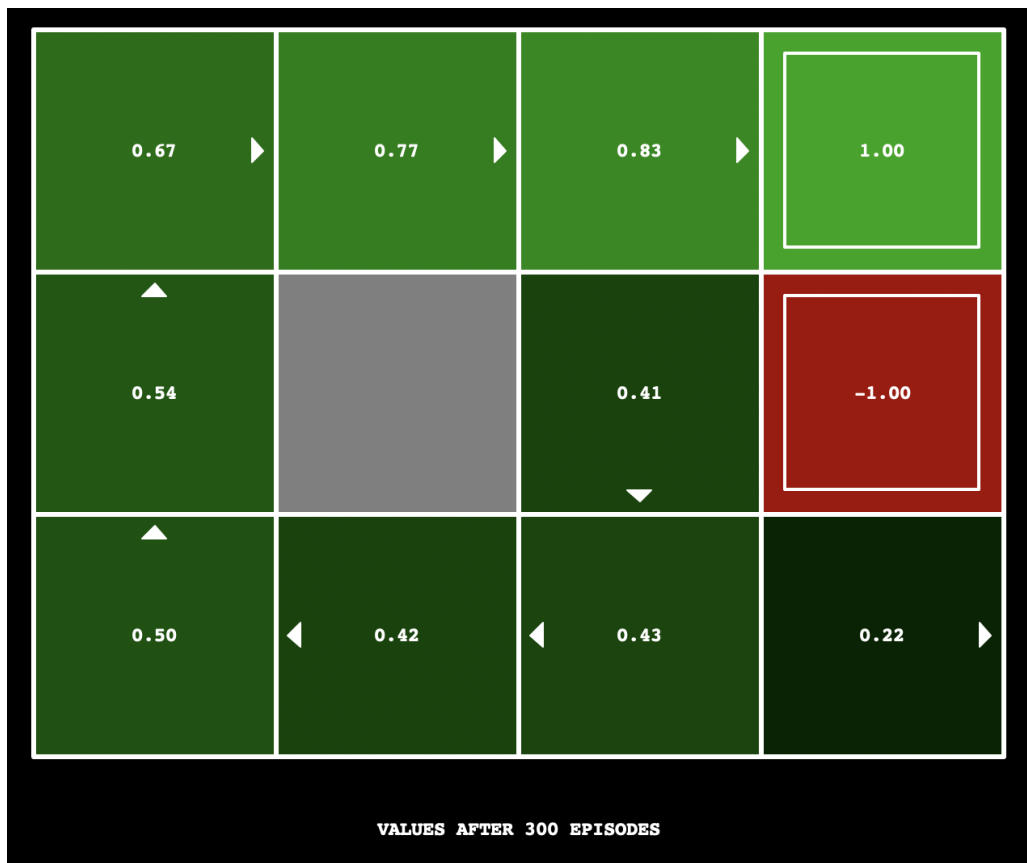


Figure 6: Value estimates obtained with Q learning on BookGrid after 300 episodes

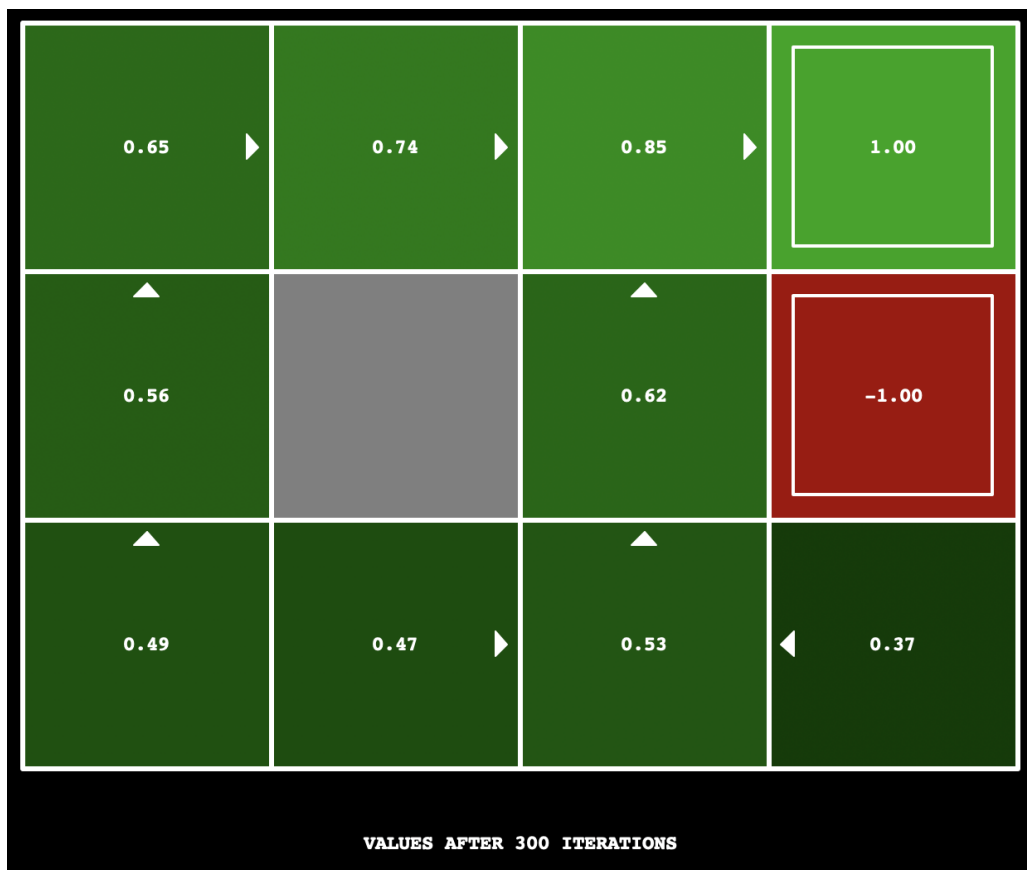


Figure 7: Value estimates obtained with value iteration on BookGrid after 300 iterations



Figure 8: Value estimates obtained with Q learning on BookGrid after 10000 episodes