

Reinforcement Learning

Lecture 4

Georg Martius

Distributed Intelligence / Autonomous Learning Group, Uni Tübingen, Germany

November 5, 2024



Model-free Control

Overview

- ▶ Last time: **unknown** MDP
 - ▶ how to estimate the value function
 - ▶ no model (model-free)
- ▶ Today: **Solve** an **unknown** MDP
 - ▶ optimize the value function
 - ▶ obtain optimal policy
 - ▶ also without model
- ▶ Next time: Bandits – a special case of RL problems with a single state and exploration (given by Claire Vernade)

Plan for today

1. Q-function
2. Monte-Carlo Control
3. Temporal Difference Learning
4. On-Policy vs. Off-Policy
5. Off-Policy Learning (Q-learning)

Where to apply model-free control?

- ▶ MDP model is **unknown**, but experience **can be sampled**

Examples:

- ▶ Robot control
- ▶ Tokamak plasma control
- ▶ Protein Folding

- ▶ MDP model is **known**, but is **too big** to use, except by samples

Examples:

- ▶ Game of Go
- ▶ Elevator optimization
- ▶ Optimizing decision making problems in simulation

What stops us from optimizing the value function?

We know how to estimate the value function: Monte-Carlo or TD

How to derive the (greedy) policy? (deterministic notation)

$$\pi'(s) = \arg \max_{a \in \mathcal{A}} \mathcal{R}_s^a + \mathcal{P}_{ss'}^a V(s')$$

- ➔ greedy policy improvement over $V(s)$ requires the MDP or a **model** of it

What to do? **Learn an action-value function!**

Action-value function (reminder)

$q_\pi(s, a)$: value of performing action a in state s and then following the policy
 $Q(s, a)$ is an estimate of $q_\pi(s, a)$

$$\pi'(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$$


- ➔ greedy policy improvement over $Q(s, a)$ is **model-free**

Policy Iteration (reminder)

Policy Iteration

Given a policy π

1. **Policy Evaluation:** Evaluate the policy π
 ➔ estimate v_π
2. **Policy Improvement:** $\pi' > \pi$
 e.g. Improve the policy by acting greedily with respect to v_π **requires model**
3. iterate until value function does not change




- ▶ Policy iteration always converges to π^*

Policy Iteration with Action-Value Function

Policy Iteration with Action-Value Function

Given a policy π

1. **Policy Evaluation:** Evaluate the policy π
 ➔ estimate **action-value function** q_π
2. **Policy Improvement:** $\pi' > \pi$
 e.g. Improve the policy by acting greedily with respect to q_π ✓
3. iterate until **action-value** function does not change



Greedy action selection



"Behind one door is tenure - behind the other is flipping burgers at McDonald's."
Copyright © 2003 David Farley, d-farley@ibiblio.org

Example: There are two doors in front of you:

- ▶ You open the left door and get reward 0
 $V(\text{left}) = 0$
- ▶ You open the right door and get reward +1
 $V(\text{right}) = +1$
- ▶ You open the right door and get reward +3
 $V(\text{right}) = +2$
- ▶ You open the right door and get reward +2
 $V(\text{right}) = +2$

Are you sure you've chosen the best door?

Exploration is needed!

Does ϵ -Greedy policy harm improvement? NO

i.e.: does Policy Iteration still converge?

Theorem: ϵ -Greedy Policy Improvement

For any ϵ -greedy policy π ,
the ϵ -greedy policy π' with respect to q_π
is an improvement, i.e. $v_{\pi'}(s) \geq v_\pi(s)$

$$\begin{aligned}
 q_\pi(s, \pi'(s)) &= \sum_{a \in \mathcal{A}} \pi'(a | s) q_\pi(s, a) \\
 &= \epsilon/m \sum_{a \in \mathcal{A}} q_\pi(s, a) + (1 - \epsilon) \max_{a \in \mathcal{A}} q_\pi(s, a) \\
 &\geq \epsilon/m \sum_{a \in \mathcal{A}} q_\pi(s, a) + (1 - \epsilon) \sum_{a \in \mathcal{A}} \frac{\pi(a | s) - \epsilon/m}{1 - \epsilon} q_\pi(s, a) \\
 &= \sum_{a \in \mathcal{A}} \pi(a | s) q_\pi(s, a) = v_\pi(s)
 \end{aligned}$$

Therefore: $v_{\pi'}(s) \geq v_\pi(s)$

ϵ -Greedy exploration

Simplest idea for ensuring continual exploration

- ▶ All m actions are tried with non-zero probability
- ▶ With probability $1 - \epsilon$ choose the greedy action
- ▶ With probability ϵ choose an action at random

$$\pi(a | s) = \begin{cases} \frac{\epsilon}{m} + 1 - \epsilon & \text{if } a^* = \arg \max_{a \in \mathcal{A}} Q(s, a) \\ \frac{\epsilon}{m} & \text{otherwise} \end{cases}$$

Monte-Carlo Action-value function estimation

1. trajectory using π : $\tau = \{S_1, A_1, R_2, \dots, S_T\} \sim \pi$

2. For each state S_t and action A_t in τ :

$$N(S_t, A_t) \leftarrow N(S_t, A_t) + 1$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{1}{N(S_t, A_t)} (G_t - Q(S_t, A_t))$$

3. iterate


We have the ingredients

- ▶ Action-value estimation
- ▶ Policy: ϵ -greedy

Monte-Carlo Control / Policy Iteration

Every episode we iterate:

1. Policy evaluation: Monte-Carlo action-value estimation $Q \approx q_\pi$
2. Policy improvement: ϵ -greedy policy improvement



Convergence?

Greedy in the Limit with Infinite Exploration (GLIE)

- ▶ All state-action pairs are explored infinitely many times,

$$\lim_{k \rightarrow \infty} N_k(s, a) = \infty$$

- ▶ The policy converges on a greedy policy,

$$\lim_{k \rightarrow \infty} \pi_k(a | s) = \mathbb{I}[a = \arg \max_{a' \in \mathcal{A}} Q_k(s, a')]$$

- ▶ Adapt exploration rate as $\epsilon \leftarrow \frac{1}{k}$
 ➔ converges to greedy policy
- ▶ infinite exploration because $\sum_{k=1}^{\infty} \frac{1}{k} = \infty$

Theorem

GLIE Monte-Carlo Control converges to the optimal action-value function, $Q(s, a) \rightarrow q^*(s, a)$

Convergence?

What do you think is needed?

1. nothing, it converges always ✗
2. it does not converge ✗
3. a lot of iterations ✓
4. every state, action pair is visited infinitely often ✓
5. a greedy policy ✗
6. policy needs to become greedy in the limit ✓
7. policy needs to keep exploring ✓

Greedy in the Limit with Infinite Exploration (GLIE)

- ▶ All state-action pairs are explored infinitely many times,

$$\lim_{k \rightarrow \infty} N_k(s, a) = \infty$$

- ▶ The policy converges on a greedy policy,

$$\lim_{k \rightarrow \infty} \pi_k(a | s) = \mathbb{I}[a = \arg \max_{a' \in \mathcal{A}} Q_k(s, a')]$$


Example: Blackjack



- ▶ Goal: get as close as possible to 21 points but not above
- ▶ Counting: Face-cards: 10, Ace: 1 or 11, other: cards their value
- ▶ start with two cards, dealer has one card open

Example: Monte-Carlo Control in Blackjack

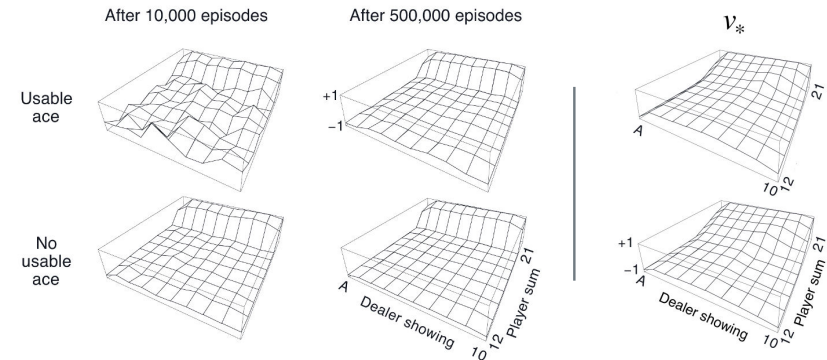
- Dealer: **stick** if sum of cards ≥ 17 , otherwise **twist**



Break

Compare: Optimal Value vs Policy Evaluation

- Policy: **stick** if sum of cards ≥ 20 , otherwise **twist** (suboptimal policy)
- Dealer: **stick** if sum of cards ≥ 17 , otherwise **twist**



Using Temporal-difference (TD) learning

What are the advantages of TD over MC?


- Lower variance
- online
- incomplete sequences
- exploits Markov property

Natural idea: use TD instead of MC

Simplest TD Q update:

$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$

S A R S A



SARSA

1. Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$
2. Repeat (for each episode):
 - 2.1 $S \leftarrow$ from environment
 - 2.2 $A \sim \pi_\epsilon^Q(S)$ (ϵ -greedy w.r.t. Q)
 - 2.3 Repeat (for each step of episode):
 - 2.3.1 Take action A , observe R, S'
 - 2.3.2 $A' \sim \pi_\epsilon^Q(S')$ (ϵ -greedy w.r.t. Q)
 - 2.3.3 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$
 - 2.3.4 $S \leftarrow S'; A \leftarrow A'$
 - 2.4 until S is terminal

Convergence of Sarsa

Theorem


Sarsa converges to the optimal action-value function, $Q(s, a) \rightarrow q^*(s, a)$, under the following conditions:

- ▶ GLIE sequence of policies $\pi_t(a | s)$ (decay of exploration rate)
- ▶ Robbins-Monro sequence of step-sizes α_t

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$


Windy Gridworld Example



Reward: $R = -1$ per timestep until reaching goal

Undiscounted: $\gamma = 1$

SARSA on the Windy Gridworld Example



SARSA(λ) Algorithm

Naturally we can combine n -step TD updates with SARSA \rightarrow SARSA(λ)
Use eligibility traces (backward view):


SARSA(λ)

1. Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$
2. Repeat (for each episode):
 - 2.1 $E(s, a) = 0, \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$
 - 2.2 $S \leftarrow$ from environment
 - 2.3 $A \sim \pi_\epsilon^Q(S)$ (ϵ -greedy w.r.t. Q)
 - 2.4 Repeat (for each step of episode):
 - 2.4.1 Take action A , observe R, S'
 - 2.4.2 $A' \sim \pi_\epsilon^Q(S')$ (ϵ -greedy w.r.t. Q)
 - 2.4.3 $\delta \leftarrow R + \gamma Q(S', A') - Q(S, A)$
 - 2.4.4 $E(S, A) \leftarrow E(S, A) + \delta$
 - 2.4.5 For all $s \in \mathcal{S}, a \in \mathcal{A}(s)$:


$$Q(s, a) \leftarrow Q(s, a) + \alpha \delta E(s, a)$$

$$E(s, a) \leftarrow \gamma \lambda E(s, a)$$
 - 2.4.6 $S'; A \leftarrow A'$
 - 2.5 until S is terminal

SARSA(λ) update on the Windy Gridworld Example



Break



On-policy vs. Off-Policy

On-policy

- ▶ “Learn while doing the job”
- ▶ Learn about policy π from experience sampled with π

Examples:

- ▶ MC Control/Policy Iteration
- ▶ SARSA

Off-policy

- ▶ “Look over someone’s shoulder”
- ▶ Learn about policy π from experience sampled with μ

Examples: (discussed next)

- ▶ Importance Sampling MC
- ▶ Q-Learning

Why is off-policy learning a good idea?


- ▶ can learn from demonstrations
- ▶ can re-use all experience (data from policies $\pi_1, \pi_2, \dots, \pi_{t-1}$)
- ▶ learn about **optimal** policy while following **exploratory** policy
- ▶ learn about **multiple** policies while following **one** policy

Importance Sampling

We get samples from current policy, but would like to compute the expectation with respect to a different policy.

Importance Sampling

$$\begin{aligned}\mathbb{E}_{X \sim P}[f(X)] &= \sum P(X) f(X) \\ &= \sum Q(X) \frac{P(X)}{Q(X)} f(X) \\ &= \mathbb{E}_{X \sim Q} \left[\frac{P(X)}{Q(X)} f(X) \right]\end{aligned}$$



⇒ estimate expectation of a different distribution

Off-policy without Importance Sampling

Off-policy learning of action-values $Q(s, a)$

- ▶ $q_\pi(s, a)$: value of doing action a in state s and then following π
- ▶ we want to approximate $q_\pi(s, a)$,
but get data from policy μ : $A_t \sim \mu(\cdot | S_t)$
- ▶ No problem!
Current action A_t comes from exploration policy μ , but then we follow π .
 $A' \sim \pi(\cdot | S_{t+1})$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left(R_{t+1} + \underbrace{\gamma Q(S_{t+1}, A')}_{\approx v_\pi(S_{t+1})} - Q(S_t, A_t) \right)$$

⇒ Q estimates $q_\pi(s, a)$

Importance Sampling for Off-Policy Monte-Carlo

Have: return G_t from policy μ

Importance Sampling: reweight according to similarity between policies

- ▶ Multiply importance sampling corrections along whole episode

$$G_t^{\pi/\mu} = \frac{\pi(A_t | S_t)}{\mu(A_t | S_t)} \frac{\pi(A_{t+1} | S_{t+1})}{\mu(A_{t+1} | S_{t+1})} \dots \frac{\pi(A_T | S_T)}{\mu(A_T | S_T)} G_t$$

- ▶ Update value towards corrected return

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^{\pi/\mu} - V(S_t))$$

- ▶ **Warning:** cannot use if μ is zero when π is non-zero
- ▶ Importance sampling can **dramatically increase variance**

Similarly possible for TD: weight TD-target by $\frac{\pi}{\mu}$

Q-Learning

Let both exploration- and target policy **improve!**


- ▶ The target policy π is **greedy** w.r.t. $Q(s, a)$ (approaching π^*)

$$\pi(S_{t+1}) = \arg \max_{a'} Q(S_{t+1}, a')$$

- ▶ The behavior policy μ is e.g. **ϵ -greedy** w.r.t. $Q(s, a)$
- ▶ The Q-learning target then simplifies:

$$\begin{aligned}& R_{t+1} + \gamma Q(S_{t+1}, A') & A' \sim \pi(\cdot | S_{t+1}) \\ &= R_{t+1} + \gamma Q\left(S_{t+1}, \arg \max_{a'} Q(S_{t+1}, a')\right) \\ &= R_{t+1} + \max_{a'} \gamma Q(S_{t+1}, a')\end{aligned}$$

Q-Learning Control update



$$Q(S, A) \leftarrow Q(S, A) + \alpha \left(R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right)$$

Theorem

Q-Learning (Control) converges to the optimal action-value function, $Q(s, a) \rightarrow q^*(s, a)$

Q-Learning Demo

```
python3 gridworld.py -a q -k 20 -n 0
```

(You can do the same after completing the homeworks)

Let's look at it after 100 or 1000 iterations:

```
python3 gridworld.py -a q -k 100 -n 0 -q
```

What happens if we add noise?


```
python3 gridworld.py -a q -k 1000 -n 0.2 -q
```

Q-Learning Algorithm

Q-Learning

1. Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$
2. Repeat (for each episode):
 - 2.1 $S \leftarrow$ from environment
 - 2.2 Repeat (for each step of episode):
 - 2.2.1 $A \sim \pi_{\epsilon}^Q(S)$ (ϵ -greedy w.r.t. Q)
 - 2.2.2 Take action A , observe R, S'
 - 2.2.3 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_{a'} Q(S', a') - Q(S, A)]$
 - 2.2.4 $S \leftarrow S'$
 - 2.3 until S is terminal

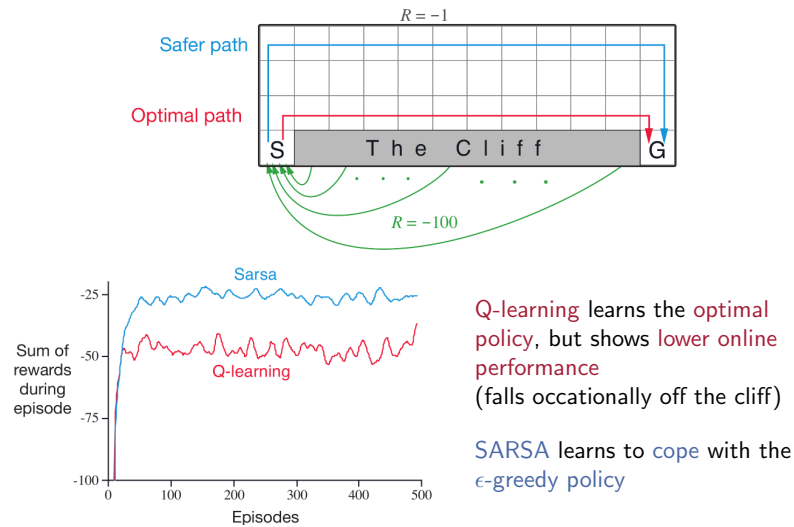
Cliff Walking Example



Fixed ϵ -greedy exploration policy. No noise in environment.

1. Which path will be learned by Q-Learning?
2. Which path will be learned by SARSA? (Poll)

Cliff Walking Example



Relationship between DP and TD (cont)


Full Backup (DP)	Sample Backup (TD)
Iterative Policy Evaluation $V(s) \leftarrow \mathbb{E}[R + \gamma V(S') s]$	TD Learning $V(S) \stackrel{\alpha}{\leftarrow} R + \gamma V(S')$
Q-Policy Iteration $Q(s, a) \leftarrow \mathbb{E}[R + \gamma Q(S', A') s, a]$	SARSA $Q(S, A) \stackrel{\alpha}{\leftarrow} R + \gamma Q(S', A')$
Q-Value Iteration $Q(s, a) \leftarrow \mathbb{E}\left[R + \gamma \max_{a' \in \mathcal{A}} Q(S', a') s, a\right]$	Q-Learning $Q(S, A) \stackrel{\alpha}{\leftarrow} R + \gamma \max_{a' \in \mathcal{A}} Q(S', a')$

where $x \stackrel{\alpha}{\leftarrow} y \equiv x \leftarrow x + \alpha(y - x)$

Relationship between DP and TD

	Full Backup (DP)	Sample Backup (TD)
Bellman Expectation Equation for $v_{\pi}(s)$	Iterative policy Evaluation $v_{\pi}(s) \leftarrow s$ $v_{\pi}(s') \leftarrow s'$	TD Learning
Bellman Expectation Equation for $q_{\pi}(s, a)$	Q-policy Iteration $q_{\pi}(s, a) \leftarrow s, a$ $q_{\pi}(s', a') \leftarrow s', a'$	SARSA
Bellman Optimality Equation for $q^*(s, a)$	Q-Value Iteration $q_*(s, a) \leftarrow s, a$ $q_*(s', a') \leftarrow s', a'$	Q-Learning

Questions?



[Image: globalrobots.com]