

# Reinforcement Learning

## Lecture 6

Georg Martius

Distributed Intelligence / Autonomous Learning Group, Uni Tübingen, Germany

November 19, 2024

EBERHARD KARLS  
UNIVERSITÄT  
TÜBINGEN



- ▶ Last time: Multiarmed Bandits
  - ▶ RL without states
  - ▶ Optimal exploration in noisy reward environments
- ▶ Today: RL with **function approximation**
  - ▶ Value function approximation
- ▶ Next time: policy gradient

# Value Function Approximation

Slides adapted from David Silver, Material from Sutton & Barto's RL book

# Plan for today

1. Introduction
2. Incremental Methods
3. Batch Methods

What is the problem with the tabular case?

- ▶ State spaces are **exponentially large** for many interesting problems
  - ▶ Backgammon:  $10^{20}$  states
  - ▶ Go (Board Game):  $10^{170}$
- ▶ Or the state space is **continuous**
  - ▶ Robot control

What to do?

- ▶ **Approximate the value function** by a function approximator.

# Value Function Approximation

- ▶ The tabular value function
  - ▶ Every state  $s$  has an entry  $V(s)$
  - ▶ Or every state-action pair  $s, a$  has an entry  $Q(s, a)$
- ▶ Problem with large MDPs:
  - ▶ There are **too many states** and/or actions to store in memory
  - ▶ It is **too slow to learn** the value of each state individually
- ▶ Solution for large MDPs:
  - ▶ Estimate value function with function approximation


$$\hat{v}(s, \mathbf{w}) \approx v_{\pi}(s)$$

$$\text{or } \hat{q}(s, a, \mathbf{w}) \approx q_{\pi}(s, a)$$


- ▶ **Generalize** from seen states to unseen states
- ▶ **Update** parameter  $\mathbf{w}$  using MC or TD learning
- ▶ Notation: approximations denoted with “hat”

# Types of Value Function Approximation


Value function



Action value function



Action value function for discrete actions



# Which Function Approximator?

There are many function approximators, for instance:

- ▶ Linear combinations of features
- ▶ Neural network
- ▶ Decision tree
- ▶ Nearest neighbors
- ▶ Fourier / wavelet bases
- ▶ ...



# Which Function Approximator?

Differentiable function approximators will be handy:

- ▶ Linear combinations of features
- ▶ Neural network
- ▶ Decision tree
- ▶ Nearest neighbors
- ▶ Fourier / wavelet bases
- ▶ ...

the value changes with policy,  
so it is a moving target (not stationary)

We require a training method that is suitable for **non-stationary, non-iid** (not identically and independently distributed) data

Suitable learning algorithm: (Stochastic) Gradient Descent

# Value Function Approximation

using supervised learning

- ▶ Value function parametrized by parameters  $\mathbf{w}$ :  $\hat{v}(s, \mathbf{w})$
- ▶ Goal: find parameter vector  $\mathbf{w}$  minimizing difference between approximate value function  $\hat{v}(s, \mathbf{w})$  and true value function  $v_\pi(s)$

$$L(\mathbf{w}) = \mathbb{E}_\pi \left[ \left( v_\pi(S) - \hat{v}(S, \mathbf{w}) \right)^2 \right]$$


- Loss function:

$$L(\mathbf{w}) = \mathbb{E}_{\pi} \left[ \left( v_{\pi}(S) - \hat{v}(S, \mathbf{w}) \right)^2 \right]$$


- Gradient descent finds a local minimum

$$\begin{aligned} \Delta \mathbf{w} &= -\frac{1}{2} \alpha \nabla_{\mathbf{w}} L(\mathbf{w}) \\ &= \alpha \mathbb{E}_{\pi} \left[ \left( v_{\pi}(S) - \hat{v}(S, \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w}) \right] \end{aligned}$$

For convex functions:



The general case:



[Image: hackernoon.com]

# Stochastic Gradient Descent (SGD)

- ▶ General loss:


$$L(\mathbf{w}) = \mathbb{E}_{(x,y) \sim \mathcal{D}} \left[ d \left( \underbrace{y}_{\text{target}}, \underbrace{\hat{y}(x)}_{\text{prediction}} \right) \right]$$

- ▶ Loss is expected empirical error: sum over examples (batch)
- ▶ SGD: update parameters on every example:

$$\Delta \mathbf{w} = -\alpha \cancel{\frac{1}{N} \sum_i} \nabla_{\mathbf{w}} d(y^{(i)}, \hat{y}^{(i)})$$

expected update is equal to full gradient update

- ▶ Minibatches:  
average gradient over a small # of examples



instead of averaging gradient over all samples,  
do it in small batches

Represent state by a feature vector

$$\mathbf{x}(S) = \begin{pmatrix} x_1(S) \\ x_2(S) \\ \vdots \\ x_n(S) \end{pmatrix}$$

For example:

- ▶ Distance of robot from landmarks
- ▶ Trends in the stock market
- ▶ Piece and pawn configurations in chess

# Linear Value Function Approximation

- ▶ Represent value function by a linear combination of features

$$\hat{v}(S, \mathbf{w}) = \mathbf{x}(S)^\top \mathbf{w} = \sum_{j=1}^n \mathbf{x}_j(S) \mathbf{w}_j$$

- ▶ Objective function is quadratic in parameters  $\mathbf{w}$

$$L(\mathbf{w}) = \mathbb{E}_\pi \left[ \left( v_\pi(S) - \mathbf{x}(S)^\top \mathbf{w} \right)^2 \right]$$

- ▶ Stochastic gradient descent converges on global optimum  
Update rule is particularly simple

$$\begin{aligned} \nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w}) &= \mathbf{x}(S) \\ \Delta \mathbf{w} &= \alpha (v_\pi(S) - \hat{v}(S, \mathbf{w})) \mathbf{x}(S) \end{aligned}$$

- ▶ Update = step-size  $\times$  prediction error  $\times$  feature value

# One-hot features = Tabular representation

Given a one-hot feature vector

$$\mathbf{x}(S) = \begin{pmatrix} \mathbb{I}[S = s_1] \\ \vdots \\ \mathbb{I}[S = s_n] \end{pmatrix}$$

if in state  $s_1$ :  
 $\mathbf{x}(S) = (1, 0, \dots, 0)$

$\mathbf{x}(S)^\top \mathbf{w}$  yields the individual “table entries”

# Iterative Prediction Algorithms

In contrast to supervised learning, we do **not have access to**  $v_{\pi}(s)$

- ▶ In RL: only rewards are given
- ▶ In practice, we substitute a target for  $v_{\pi}(s)$ 
  - ▶ For MC, use the return  $G_t$

$$\Delta \mathbf{w} = \alpha (G_t - \hat{v}(S_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S_t, \mathbf{w})$$

- ▶ For TD(0): use TD-target  $R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w})$

$$\Delta \mathbf{w} = \alpha (R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S_t, \mathbf{w})$$

- ▶ For TD( $\lambda$ ): use  $\lambda$ -return  $G_t^{\lambda}$

$$\Delta \mathbf{w} = \alpha (G_t^{\lambda} - \hat{v}(S_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S_t, \mathbf{w})$$

- ▶ as before the TD-target is biased (we use our own estimate)
- ▶ Note: For TD: target is not independent of  $\mathbf{w}$ .  
**Semi-gradient methods**: do not necessarily converge



# Iterative Prediction Algorithm

## for Linear Features

Consider the case of **linear features**:  $\hat{v}(s, \mathbf{w}) = \mathbf{x}(s)^\top \mathbf{w}$ :

- ▶ The derivative  $\nabla_{\mathbf{w}} \hat{v}(s, \mathbf{w})$  is simply  $\mathbf{x}(s)$
- ▶ MC:

$$\begin{aligned}\Delta \mathbf{w} &= \alpha (G_t - \hat{v}(S_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S_t, \mathbf{w}) \\ &= \alpha (G_t - \hat{v}(S_t, \mathbf{w})) \mathbf{x}(S_t)\end{aligned}$$

- ▶ TD(0):

$$\begin{aligned}\Delta \mathbf{w} &= \alpha \underbrace{(R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w}))}_{\delta} \nabla_{\mathbf{w}} \hat{v}(S_t, \mathbf{w}) \\ &= \alpha \delta_t \mathbf{x}(S_t)\end{aligned}$$


- ▶ TD( $\lambda$ ): Forward view:  $\Delta \mathbf{w} = \alpha (G_t^\lambda - \hat{v}(S_t, \mathbf{w})) \mathbf{x}(S_t)$
- ▶ TD( $\lambda$ ): Backward view with Eligibility trace  $E$ :

$$\begin{aligned}E_t &= \gamma \lambda E_{t-1} + \mathbf{x}(S_t) \\ \Delta \mathbf{w} &= \alpha \delta_t E_t\end{aligned}$$

# For control: Action-Value Function

We can do the same procedure for the **action-value** function.

$$\hat{q}(s, a, \mathbf{w}) \approx q_{\pi}(s, a)$$



- Policy evaluation: update  $\mathbf{w}$
- Policy improvement:  $\epsilon$ -greedy policy improvement

# Linear Action-Value Function Approximation

- Represent state and action by a feature vector

$$\mathbf{x}(S, A) = \begin{pmatrix} \mathbf{x}_1(S, A) \\ \vdots \\ \mathbf{x}_n(S, A) \end{pmatrix}$$

- Represent action-value function by linear combination of features:

$$\hat{q}(S, A, \mathbf{w}) = \mathbf{x}(S, A)^\top \mathbf{w} = \sum_{j=1}^n \mathbf{x}_j(S, A) \mathbf{w}_j$$

- Stochastic gradient descent update

$$\begin{aligned} \nabla_{\mathbf{w}} \hat{q}(S, A, \mathbf{w}) &= \mathbf{x}(S, A) \\ \Delta \mathbf{w} &= \alpha \left( \underbrace{q_\pi(S, A)}_{\text{target}} - \hat{q}(S, A, \mathbf{w}) \right) \mathbf{x}(S, A) \end{aligned}$$

- Replace **target** by:  $G_t$  or TD-target ...

SARSA: using TD(0) target (semi-gradient method)

## SARSA with function approximators (episodic)

Given:  $\hat{q}(s, a, \mathbf{w}) : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^d \rightarrow \mathbb{R}$

1. Initialize  $\mathbf{w}$

2. Repeat (for each episode):

2.1  $S \leftarrow$  from environment

2.2  $A \sim \pi^{\hat{q}}(S)$  (e.g.  $\epsilon$ -greedy w.r.t.  $\hat{q}$ )

2.3 Repeat (for each step of episode):

2.3.1 Take action  $A$ , observe  $R, S'$

2.3.2 if  $S'$  terminal

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha[R - \hat{q}(S, A, \mathbf{w})] \nabla_{\mathbf{w}} \hat{q}(S, A, \mathbf{w})$$

Go to next episode

2.3.3  $A' \sim \pi^{\hat{q}}(S')$  (e.g.  $\epsilon$ -greedy w.r.t.  $\hat{q}$ )


2.3.4  $\mathbf{w} \leftarrow \mathbf{w} + \alpha \underbrace{[R + \gamma \hat{q}(S', A', \mathbf{w}) - \hat{q}(S, A, \mathbf{w})]}_{\text{TD-target}} \nabla_{\mathbf{w}} \hat{q}(S, A, \mathbf{w})$

2.3.5  $S \leftarrow S'; A \leftarrow A'$

## Topical Break / Discussion

What features can you think of for:

- ▶ a cart-pendulum



- ▶ large chain  
(random walk example with 1000 states)
- ▶ grid worlds

# Feature representations

## ► Polynomials:

- when states are numbers: positions/velocities of robot, #cars in parking, ...
- for continuous state-space with  $s_1, s_2 \in \mathbb{R}$ , consider concrete example:


$$\mathbf{x}(s) = (s_1, s_2)^\top$$

Why is that not so good?


- if  $s_1 = s_2 = 0$  then value is 0, no interaction between  $s_1$  and  $s_2$ , very inflexible
- $\mathbf{x}(s) = (1, s_1, s_2, s_1 s_2)^\top$  or  $\mathbf{x}(s) = (1, s_1, s_2, s_1 s_2, s_1^2, s_2^2, s_1^2 s_2, s_1 s_2^2, s_1^2 s_2^2)^\top$
- for degree  $n$  and state-space  $k$  the number of features grows exponentially in  $k$

## ► Fourier Basis

- Cosine basis:  $\mathbf{x}_i(s) = \cos(i\pi s), s \in [0, 1]$




- Value Error (RMSE) for 1000 random states example and gradient MC  $\Rightarrow$




# Feature representations (Cont)


- Coarse Coding: inside circle: feature = 1



Narrow generalization




Broad generalization




Asymmetric generalization

- Initial generalization:



- ▶ Radial basis functions (RBF)

$$\mathbf{x}_i(\mathbf{s}) = \exp(-\|\mathbf{s} - \mathbf{c}_i\|^2 / 2\sigma_i^2)$$




- ▶ How to choose  $c_i$ 's and  $\sigma_i$ 's?
- ▶ Scales badly with the dimensionality




# Non-linear function approximators

## Neural Networks:



they are doing function approximation and regression at the same time!  
No need to do feature engineering

Direct approximation (implicit feature learning):









# Example: Linear SARSA with Coarse Coding

## Mountain Car Environment

MOUNTAIN CAR

Goal



# Off-policy and Convergence

- ▶ Exactly the same can be done for the off-policy case (Q-Learning)
- ▶ In the tabular case we have proven TD to converge
- ▶ With function approximation and off-policy this cannot be done

## Simple Example

- ▶ Two states: with features  $w$  and  $2w$ :




- ▶ Assume initial  $w = 10$
- ▶ Transition in data: left to right
- ▶  $v(l) = 10$  to  $v(r) = 20$ : TD error = 10
- ▶ Update for  $w$  with  $\alpha = 0.1$ :  $w = 11$
- ▶ ➡  $v(r) = 22$
- ▶ Next update TD error = 11 (error is even larger!)


Incomplete MDP, so maybe not convincing

error is getting  
larger, but maybe  
not best of features

## Baird's Counterexample



## Semi-gradient Off-policy TD



TD off-policy diverges

Features / representation needs to be **adapted** and should **not be fixed**.

# Convergence of Prediction Algorithms

	Algorithm	Table Lookup	Linear	Non-linear
On-Policy	MC	✓	✓	✓
	TD	✓	✓	✗
	Gradient TD	✓	✓	✓
Off-Policy	MC	✓	✓	✓
	TD	✓	✗	✗
	Gradient TD	✓	✓	✓

- ▶ TD does not follow the gradient of any objective function, **divergence possible**
- ▶ Gradient TD follows true gradient of projected Bellman error [Maei et al, 2009]

# Convergence of Control Algorithms

Algorithm	Table Lookup	Linear	Non-linear
Monte-Carlo Control	✓	(✓)	✗
SARSA	✓	(✓)	✗
Q-Learning	✓	✗	✗
Gradient Q-Learning	✓	✓	✗

(✓) means that it fluctuate around near-optimal solution

For control with **non-linear features**, we **cannot guarantee convergence**.

# Incremental or Batch?

- ▶ So far: Incremental update, i.e. update parameters towards targets locally by SGD
- ▶ Alternative: find a parameters that fit best the experience / training data
  - ➡ Batch methods

# Least Squares Prediction

- ▶ Value function approximation  $\hat{v}(s, \mathbf{w}) \approx v_\pi(s)$
- ▶ Experience  $\mathcal{D}$ : pairs of (state, value)

$$\mathcal{D} = \{(s_1, v_1^\pi), (s_2, v_2^\pi), \dots, (s_T, v_T^\pi)\}$$

- ▶ Which parameters  $\mathbf{w}$  give the best fitting value function  $\hat{v}(s, \mathbf{w})$ ?
- ▶ Least squares algorithms:

$$\begin{aligned}\text{LS}(\mathbf{w}) &= \sum_{t=1}^T (v_t^\pi - \hat{v}(s_t, \mathbf{w}))^2 \\ &= \mathbb{E}_{\mathcal{D}} \left[ (v^\pi - \hat{v}(s, \mathbf{w}))^2 \right]\end{aligned}$$

- ▶ If  $\hat{v}$  linear in features, then closed form solution



# SGD with Experience Replay

What about using SGD, but train on data many times?

Use experience  $\mathcal{D}$ :

$$\mathcal{D} = \{(s_1, v_1^\pi), (s_2, v_2^\pi), \dots, (s_T, v_T^\pi)\}$$

Repeat:

1. Sample state value pair (or mini-batch)

$$(s, v^{\text{target}}) \sim \mathcal{D}$$

2. Apply stochastic gradient descent update:

$$\Delta \mathbf{w} = \alpha (v^{\text{target}} - \hat{v}(s, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(s, \mathbf{w})$$

Converges to least square solution

$$\mathbf{w}^\pi = \arg \min_{\mathbf{w}} \text{LS}(\mathbf{w})$$

- ▶ We need a lot of data for high-capacity models
- ▶ In the on-policy case we have to throw away all the collected experience :-)
- ▶ Off-policy should be more efficient
  - ▶ Wait: don't we have convergence problems?
  - ▶ Neural networks are very flexible and can in practice often break the parameter-state entanglement leading to divergence
  - ▶ If behavior policy is closed to target policy (e.g.  $\epsilon$ -greedy version) less of a problem
  - ▶ We can freeze the value-target, more below

# Deep Q-Networks (DQN)

## Deep Q-Networks

- ▶ Q-Learning with Deep network as function approximator
- ▶ + Experience Replay
- ▶ + Target Network

### Experience Replay

Data:  $\mathcal{D} = \{(s_t, a_t, r_{t+1}, s_{t+1})\}$

Typically a fixed maximal size  $|\mathcal{D}|$   
(old data gets deleted)

### Target Network

Two sets of parameters:  $\mathbf{w}$  and  $\mathbf{w}^{\text{target}}$

Idea: keep  $\mathbf{w}^{\text{target}}$  fixed for some time  
(old value of  $\mathbf{w}$ )

TD: error is computed using  $\mathbf{w}^{\text{target}}$

- ▶ avoids instabilities
- ▶ for a fixed  $\mathbf{w}^{\text{target}}$  it is a gradient method

DQN [Mnih et al 2015], Earlier work: Neural fitted Q [Riedmiller 2005]

## DQN

Given: Network  $\hat{q}(s, a; \mathbf{w})$ , replay buffer  $\mathcal{D}$

1. Initialize  $\mathbf{w}$ ,  $\mathbf{w}^{\text{target}} \leftarrow \mathbf{w}$
2.  $s \leftarrow$  from environment
3.  $a \sim \pi^{\hat{q}}(s)$  (e.g.  $\epsilon$ -greedy w.r.t.  $\hat{q}$ )
4. Repeat (for each step of episode):
  - 4.1 Take action  $a$ , observe  $r, s'$
  - 4.2  $n \leftarrow \begin{cases} 0 & s' \text{ is terminal} \\ 1 & \text{otherwise} \end{cases}$  (*not done* flag)
  - 4.3 Store  $(s, a, r, s', n)$  in  $\mathcal{D}$
  - 4.4 Sample random minibatch  $D \subset \mathcal{D}$
  - 4.5 Optimize with respect to:

$$\mathcal{L}(\mathbf{w}) = \mathbb{E}_{s, a, r, s', n \sim \mathcal{D}} \left[ \left( r + n \gamma \max_{a'} \hat{q}(s', a'; \mathbf{w}^{\text{target}}) - \hat{q}(s, a; \mathbf{w}) \right)^2 \right]$$

with favorit optimizer (ADAM, SGD)

5. after many update steps:  $\mathbf{w}^{\text{target}} \leftarrow \mathbf{w}$

DQN [Mnih et al 2015], Earlier work: Neural fitted Q [Riedmiller 2005]

# Every iteration or grouped training?

- ▶ Due to random sampling, the currently collected data-point has very small influence
- ▶ Often easier to collect a bunch of data (e.g. one episode)
- ▶ make a number of training steps and continue

What made Deep RL famous in 2015...




[www.youtube.com/watch?v=rQIShnTz1kU](http://www.youtube.com/watch?v=rQIShnTz1kU)


# DQN in Atari

What made Deep RL famous in 2015...

- ▶ Observation space: raw pixels ( $84 \times 84$ ), 4 frames
- ▶ End-to-end learning of Q-values from pixels  $s$
- ▶ Action space: 18 discrete joystick positions and buttons
- ▶ Network  $q(s, \mathbf{w}) \mapsto \mathbb{R}^{18}$ : direct prediction of Q-values for **all** actions
- ▶ Reward: change in score for that step



# DQN Results in Atari



Best linear learner: using the best results with handcrafted features



# Ablation

## How important is Replay and Target Network?

Scores on some Atari games				
	Replay Target-Q	Replay Q-learning	No replay Target-Q	No replay Q-learning
Breakout	316.81	240.73	10.16	3.17
Enduro	1006.30	831.25	141.89	29.10
Space Invaders	1088.94	826.33	373.22	301.99
Seaquest	2894.40	822.55	1003.00	275.81
River Raid	7446.62	4102.81	2867.66	1453.02

# Linear Least Squares Prediction and Control

- ▶ Experience replay finds least squares solution but it may take many iterations
- ▶ Using linear value function approximation  $\hat{v}(s, \mathbf{w}) = x(s)^\top \mathbf{w}$  solve the least squares solution in closed form

## Refresher: Linear Regression

- ▶ At minimum of  $LS(\mathbf{w})$ : Derivative is zero.

$$\nabla_{\mathbf{w}} L(\mathbf{w}) = 0$$

$$\sum_{t=1}^T \mathbf{x}(s_t) \left( v_t^\pi - \mathbf{x}(s_t)^\top \mathbf{w} \right) = 0$$

$$\sum_{t=1}^T \mathbf{x}(s_t) v_t^\pi = \sum_{t=1}^T \mathbf{x}(s_t) \mathbf{x}(s_t)^\top \mathbf{w}$$

$$\mathbf{w} = \left( \sum_{t=1}^T \mathbf{x}(s_t) \mathbf{x}(s_t)^\top \right)^{-1} \sum_{t=1}^T \mathbf{x}(s_t) v_t^\pi$$

- ▶ For  $N$  features, direct solution time is  $O(N^3)$
- ▶ Incremental solution time is  $O(N^2)$  using Sherman-Morrison


# Convergence of Least Squares Prediction Algorithms

- ▶ As usual: plugging in different targets:  $LS\{MC|TD|TD(\lambda)\}$

	Algorithm	Table Lookup	Linear	Non-linear
On-Policy	MC	✓	✓	✓
	LSMC	✓	✓	–
	TD	✓	✓	✗
	LSTD	✓	✓	–
Off-Policy	MC	✓	✓	✓
	LSMC	✓	✓	–
	TD	✓	✗	✗
	LSTD	✓	✓	–

- ▶ On-Policy Least Squares needs lots of samples from current policy
- ▶ Putting it inside control: Policy iteration: **LSPI**
- ▶ **LSPI** also converges in off-policy setting  
(naturally limited to the linear case)

# Questions?



[Image: [globalrobots.com](http://globalrobots.com)]